



# Streaming Trend Discovery: real-time discovery in a sea of events

Scott Haines, Twilio



**#DDSAIS16** 



## **Agenda**

- Intro and use cases for Spark at Twilio (5 min)
- Elements of a Streaming Discovery Engine (5 min)
- Principals of Valuable Data Sets (4 min)
- Refresher on Histograms and Percentiles (2 min)
- Building a Discovery Engine on Spark: Part One (15 min)
- Break?: 10 min
- Building a Discovery Engine on Spark: Part Two (15 min)
- Lessons learned (5 min)
- Questions (5 min)

Project Files on Github @ https://bit.ly/2ISI3BS





#### **Scott Haines**

#### Principal Engineer, Voice Insights @ Twilio

- 6+ years working with Streaming Systems
  - Yahoo! Games
    - Recommendations, Personalization
  - Daily Fantasy Growth / Analytics
    - Mobile Campaigns, Fraud, Trends
  - Flurry Analytics
    - Statistics based notification system
- Stats / Data Science Enthusiast
- Drummer / Snowboarder / Hot Sauce Maker / Teacher







#### **Twilio**

Where I work

We are a public, large-scale, globally available & distributed communications as-a-service company.

- Founded in 2008. IPO'd in 2016
- Telecommunications is our bread and butter
- Along with Video / SMS / MMS / Wireless Sim / Authy / Sync
- Find out more and sign up @ twilio.com

<u>twilio.com</u>



## Use cases for Spark @Twilio

- Real-Time Dashboards and Debugging
- Fraud Detection
- Carrier Behavior and Network Anomaly Detection
- Streaming ETL, Metrics and Aggregations
- ASR and Sentiment



### Voice Insights @Twilio

- Processing >3 TB /d
- >4 billion raw events /d
- >2 billion metric & sessions aggregations /d
- Powers Voice Insights / Carrier Ops / Support and more

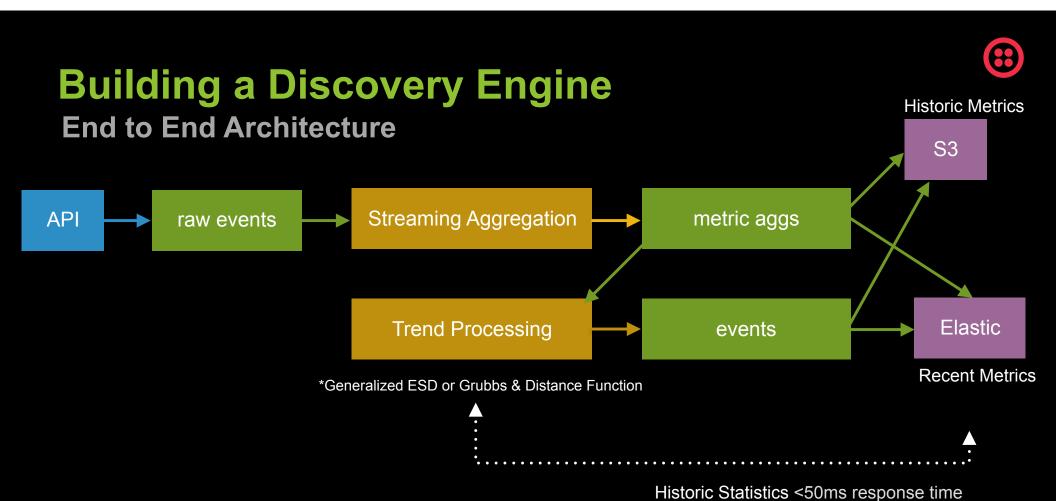




#### Next up

Elements of the Streaming Trend Discovery Architecture







#DDSAIS16

**Github** https://bit.ly/2ISI3BS



#### **Core Principals of Valuable Data Sets**

Data must be Structured. No JSON!





#### **Core Principals of Valuable Data Sets**

- Data must be Structured. No JSON!
- JSON ~= Object





#### **Core Principals of Valuable Data Sets**

- Data must be Structured. No JSON!
- JSON ~= Object
- Avro, Protobuf, Parquet are widely used and can support backwards compatible updates.





#### **Core Principals Establish a Data Contract**

Structured Data has Types and are part of the Data Contract

- Properties of the Data must be documented.
- Data must be Accurate and Complete\*
- Data must support backwards and forwards compatibility.
- Data should be easily flattened.

```
message Stats {
    optional double min
                             = 1;
    optional double p25
                             = 2;
    optional double median
                             = 3:
    optional double p75
                             = 4;
    optional double p90
                             = 5;
    optional double p95
                             = 6:
    optional double p99
                             = 7:
    optional double max
                             = 8:
    optional double mean
                             = 9:
    optional double sd
                             = 10;
    optional double variance = 11;
```





#### **Core Principals Final Thoughts**

- Structured Data allows for compile time guarantees
- Write once. Compile across languages
- Non-Java Object Serialization
- External Serialization simplifies deploys given you can simply use the last stored session state\*



<sup>\*</sup> don't corrupt your stored session data



#### **Building our mental model**

Quick Refresher on Percentiles and Histograms





Primer: What is a Percentile?

 Percentiles group numeric data in a set of observations.

$$S = \{1,2,3,4,5,6,7,8,9,10\}$$





**Primer: What is a Percentile?** 

- Percentiles group numeric data in a set of observations.
- Common splits are min, 25%,50% (median),75% and max

min:1

p25: 3.25

median: 5.50

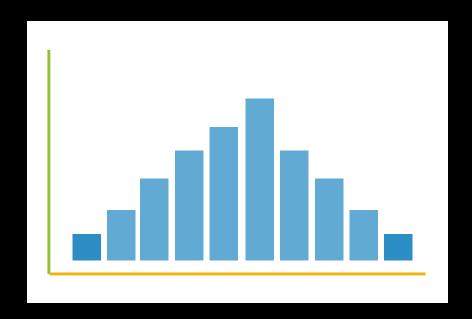
p75: 7.75 max: 10.0





Primer: What is a Histogram?

Histograms are bucketed distributions of numeric datasets

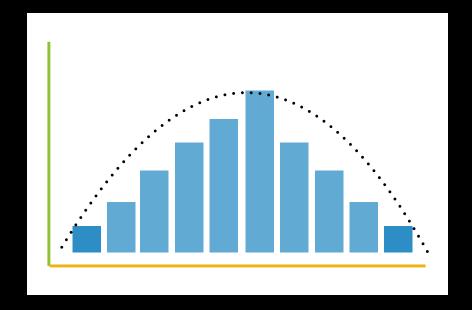






**Primer: What is a Histogram?** 

- Histograms are bucketed distributions of numeric datasets
- In a Normal Distribution, we talk about the Area Under the Curve

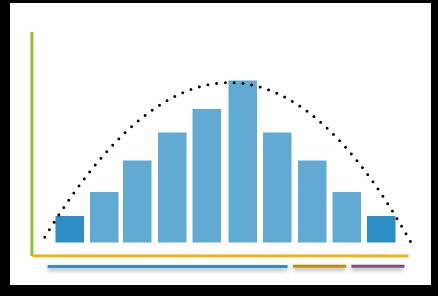






Primer: What is a Histogram?

- Histograms are bucketed distributions of numeric datasets
- A Normal Distribution, fit into the Area Under the Curve
- Histograms allow us to understand the shape of a distribution



68-95-99.7 Rule

- 68%@1sd->mean,95@2x,99.7@3x)

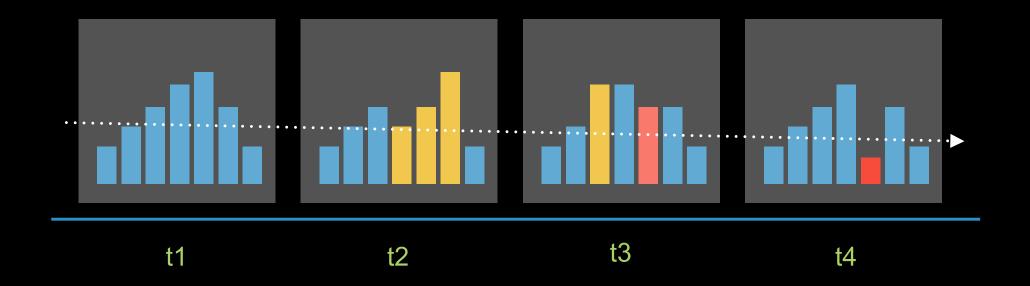


#DDSAIS16

Github https://bit.ly/2ISI3BS



**Goal: Time Series Histogram Analysis** 





#DDSAIS16

Github https://bit.ly/2ISI3BS



Deep Dive: Quantiles / Histograms in Spark

Limitations: Not easily Aggregated or hard to extract





**Deep Dive: Aggregation Extension Options** 

```
abstract class InsightsAggregation extends UserDefinedAggregateFunction {
  type T<: Product
  val log: Logger
  def productEncoder: Encoder[T]
  def productSchema: StructType = productEncoder.schema
  lazy val fieldIndex: Map[String, Int] = productSchema.fieldNames.map { fieldName =>
   fieldName -> productSchema.fieldIndex(fieldName)
 }.toMap
  val fieldLength: Int
 override def inputSchema: StructType = InsightsAggregation.inputStringRow
  override def bufferSchema: StructType = productSchema
  override def dataType: DataType = productSchema
  override def deterministic: Boolean = true
  override def initialize(buffer: MutableAggregationBuffer): Unit = {
   (0 to fieldLength).foreach { num =>
     buffer(num) = 0L
   }
 }
 override def merge(cache: MutableAggregationBuffer, state: Row): Unit = {
   (0 to fieldLength).foreach { num =>
      cache(num) = cache.getLong(num) + state.getLong(num)
   }
```

## Define a new UDF based off of the <a href="MutableAggregationBuffer">MutableAggregationBuffer</a>'s in Spark

Github https://bit.ly/2ISI3BS

**Deep Dive: Aggregation Extension Options** 

 Upside: Things are human readable



**Deep Dive: Aggregation Extension Options** 

- Upside: Things are human readable
- Downside: Performance\* / Control\*





**Deep Dive: Aggregation Extension Options: Final Thoughts** 

- Ditch the DataFrame's unstructured API and move over to Datasets typed API.
- This enables the use of flatMapGroupsWithState
- Now with a little effort you have full control over the aggregation strategy





#### Next up: after the break

Building a Discovery Engine: Part 2

- 1. Define the Problem (1m)
- 2. Build a plan of Attack (1m)
- 3. Model the Data (2m)
- 5. Build the Streaming Discovery Engine (~20m)
- 6. Hopefully time for lessons learned and Questions!







#DDSAIS16

#### Simply Connecting two endpoints is ripe for failure.







# Given Networks can go down at anytime.

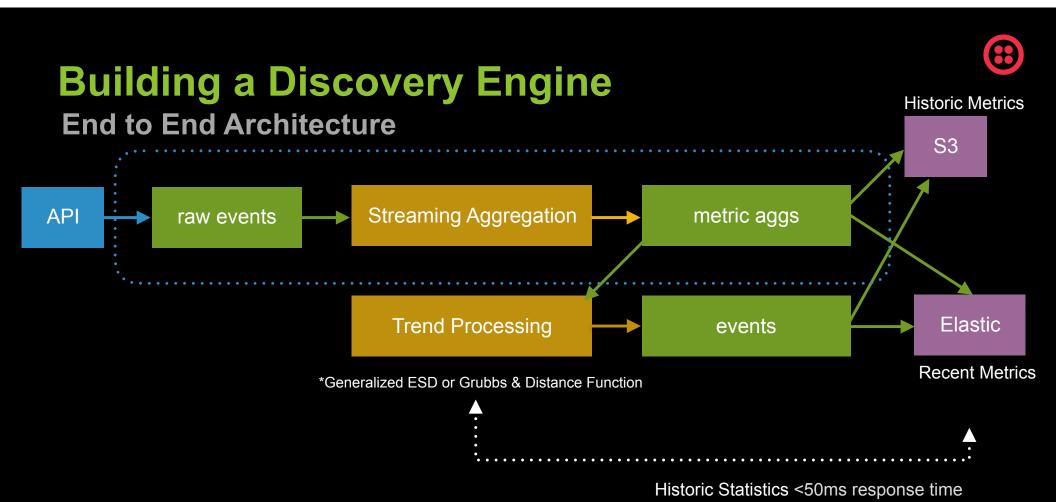


Github https://bit.ly/2ISI3BS



We should build a system to auto-discover trends in (~real-time) for our KPMs and automatically report them...









# Building a Discovery Engine Define the raw event

 Must provide the information necessary for aggregation

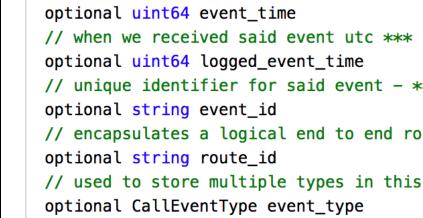
```
message CallEvent {
    // when the event occured utc
    optional uint64 event_time
    // when we received said event utc ***
    optional uint64 logged_event_time
    // unique identifier for said event - *
    optional string event_id
    // encapsulates a logical end to end ro
    optional string route_id
    // used to store multiple types in this
    optional CallEventType event_type
    // when event type is signaling_event t
    optional SignalingEvent signaling_event
    // use with the event_type and logical optional Dimensions event_dimensions
}
```





# Building a Discovery Engine Define the raw event

- Must provide the information necessary for aggregation
- id must be idempotent. Brownie points if it can be reused in Hash Partitioner



// when event type is signaling\_event t
optional SignalingEvent signaling\_event
// use with the event\_type and logical

optional Dimensions event dimensions

// when the event occured utc

message CallEvent {





## Building a Discovery Engine Define the raw event

- Must provide the information necessary for aggregation
- id must be idempotent. \*dedup key - Brownie points if it can be reused in Hash Partitioner in spark and kafka!
- event\_dimensions + event\_group can create great grouping hash

```
message CallEvent {
    // when the event occured utc
    optional uint64 event_time
    // when we received said event utc ***
    optional uint64 logged_event_time
    // unique identifier for said event - *
    optional string event_id
    // encapsulates a logical end to end ro
    optional string route_id
    // used to store multiple types in this
    optional CallEventType event_type
    // when event type is signaling_event t
    optional SignalingEvent signaling_event
    // use with the event_type and logical optional Dimensions event_dimensions
}
```





#### Define the raw event

 Allows for event evolution while supporting backwards compatibility.

```
message SignalingEvent {
  optional string name
                                          = 1;
  optional SignalingEventType event_type
                                         = 2;
  optional PddEvent pdd
                                          = 3;
  optional CallStateEvent call state
                                          = 4;
enum SignalingEventType {
  unknown_event_type = 0;
  pdd
                      = 1;
  call_state
                      = 2;
message PddEvent {
  optional float pdd = 1;
```



#DDSAIS16



#### Define the raw event

- Allows for event evolution while supporting backwards compatibility.
- Enum ordinals can be used in ML jobs without the need for string encoding

```
message SignalingEvent {
  optional string name
                                          = 1;
  optional SignalingEventType event_type
                                          = 2;
  optional PddEvent pdd
                                          = 3;
  optional CallStateEvent call state
                                          = 4;
enum SignalingEventType {
  unknown_event_type = 0;
  pdd
                      = 1;
  call_state
                      = 2;
message PddEvent {
  optional float pdd = 1;
```





#### Define the aggregation

- Includes logical time series Window
- Stores all Dimensional and Statistic
   Information

```
message MetricAggregation {
  optional string metric
                                       = 1;
  optional uint64 window_start
                                       = 2;
  optional uint64 window_end
                                       = 3;
  optional string window_interval
                                       = 4;
  optional uint32 samples
                                       = 5;
  optional Stats stats
                                       = 6;
  optional Histogram histogram
                                       = 7;
  optional Dimensions dimensions
                                       = 8:
  optional string dimension_hash
                                       = 10;
```



#DDSAIS16

**Github** https://bit.ly/2ISI3BS



Define the aggregation

Includes logical time series Window

Stores all Dimensional and Statistic
 Information

Can be used together as a strong
 Binary Group Key

```
message MetricAggregation {
  optional string metric
                                        = 1;
  optional uint64 window start
                                        = 2;
  optional uint64 window end
                                        = 3;
  optional string window interval
                                        = 4;
  optional uint32 samples
                                        = 5;
  optional Stats stats
                                        = 6;
  optional Histogram histogram
                                        = 7;
  optional Dimensions dimensions
                                        = 8;
  optional string dimension hash
                                        = 10;
```



#DDSAIS16



#### Why hash the dimensions?

- Hashing the dimensions of the aggregation simplifies the groupBy or groupByKey mechanics.
- Use guava since it ships with Spark and works well at scale for Hashing.

```
def generateId(bytes: Array[Byte]): String = {
   val hf = Hashing.murmur3_128()
   val hc = hf.hashBytes(bytes)
   BaseEncoding.base64Url()
    .omitPadding()
   .encode(hc.asBytes())
}
```





## Next up

- Putting together what we've learned with Spark
- Available in the content @ <a href="https://bit.ly/2ISI3BS">https://bit.ly/2ISI3BS</a>
- Walk through of critical components of the Application



```
override def run(): Unit = {
  // creates a streaming query that will aggregate events from kafka
  // and produce statistical histogram aggregations to be consumed downstream
  EventAggregation(config)
    .process(readKafkaStream())(spark)
    .writeStream
    .queryName("streaming.trend.discovery")
    .outputMode(eventProcessor.outputMode)
    .trigger(Trigger.ProcessingTime(Duration(config.triggerInterval)))
    .format("kafka")
    .option("topic", "spark.summit.call.aggregations")
    .options(Map(
      "checkpointLocation" -> config.checkpointPath,
      "kafka.producer.config1" -> "",
      "kafka.producer.config2" -> ""
    .start()
```

Creates the pipeline

- \* process(KDSR), inject SparkSession (implicit\*)
- \* write back to Kafka
- \* checkpoint, etc





- 1. Transform (protobuf -
- > Product)
- 2. Actual Aggregation
- \* groupedBy BinaryType
- \* fMGwS keyed by BinaryType
- \* EventTimeTimeout\*

```
def process(df: DataFrame)(implicit spark: SparkSession): Dataset[Metrics.MetricAggregation] = {
  import spark.implicits.
  import EventAggregation.transforms._
  import EventAggregation.implicits._
  implicit val metricAggsEncoder: ExpressionEncoder[MetricAggregation] = EventAggregation.implicits.aggreationEnco
  val watermarkInterval = config.watermarkIntervalMinutes
  // 1. Convert from the kafka data frame to an instance of our Metric
  val groupedMetrics = df.mapPartitions(
    _.flatMap { kd =>
      bytesToMetricRow(kd.getAs[Array[Byte]]("value"), windowInterval)
  )(metricEncoder)
  // 2.
  logger.info(s"aggregation outputMode=$outputMode windowIntervalMinutes=${config.windowIntervalMinutes} watermark
  groupedMetrics
    .withWatermark("timestamp", s"${config.watermarkIntervalMinutes} minutes")
    .groupByKey(_.metricAggregation)
    .flatMapGroupsWithState[Array[Byte], MetricAggregation](outputMode, GroupStateTimeout.EventTimeTimeout())(metr
```



def metricAggregator(

metricAggregation: Array[Byte],

metrics: Iterator[ProductMetric],

state: GroupState[Array[Byte]]): Iterator[MetricAggregation] = {



- StateSpec takes a StateFunc
- This is for Key/Value pair grouped data
- Key is [A]. It can be Binary.
   Cool
- Take an Iterator[B]
- Stores a GroupState[C]
- Emits Iterator[D]



#DDSAIS16

Github https://bit.ly/2ISI3BS



metricAggregation

```
val metricData = state.getOption match {
  case Some(lastState) =>
    logger.info(s"recovered the last state")
   // union the last and new sketches
    val lastSketch = DoubleSketch(lastState)
   val currentSketch = DoubleSketch.sketch(k, metrics.map { m => m.value }.toArray)
   val unionSketch = DoubleSketch.union(k, lastSketch, currentSketch)
   // store updated metrics
   logger.info(s"state=merge num.records=${unionSketch.getN}")
   unionSketch.toByteArray(true)
  case None =>
   val s = DoublesSketch.builder().setK(k).build()
    metrics.foreach { m =>
      s.update(m.value)
    logger.info(s"state=new num.records=${s.getN}")
    s.toByteArray(true)
state.update(metricData)
val updateTimeout = aggregation.getWindowEnd
logger.info(s"updating.timeout.timestamp.to=$updateTimeout")
state.setTimeoutTimestamp(updateTimeout)
```





- <u>Data Sketches</u> from Yahoo\* provides serializable Quantile Sketches with Histogram capabilities.
- Data Sketches are a light weight representation of a complete set of data (on a diet - think HyperLogLog)
- We use this on our team to enable performant aggregation of complicated statistics without storing all rows in Spark memory.

```
object DoubleSketch {
  val k: Int = 512
  def apply(serialized: Array[Byte]): DoublesSketch = DoublesSketch.wrap(Memory.wrap(serialized))
  def sketch(data: Array[Double]): DoublesSketch = sketch(k, data)
  def sketch(k: Int, data: Array[Double]): DoublesSketch = {
    val s = DoublesSketch.builder().setK(k).build()
   data.foreach { s.update }
  }
  def union(k: Int, a: Array[Double], b: Array[Double]): DoublesSketch = {
    val u = DoublesUnion.builder().setMaxK(k).build()
    u.update(sketch(k, a))
    u.update(sketch(k, b))
    u.getResult
  def union(k: Int, a: DoublesSketch, b: DoublesSketch): DoublesSketch = {
    val u = DoublesUnion.builder().setMaxK(k).build()
    u.update(a)
    u.update(b)
    u.getResult
```





- <u>Data Sketches</u> from Yahoo\* provides serializable Quantile Sketches with Histogram capabilities.
- Data Sketches are a light weight representation of a complete set of data (on a diet - think HyperLogLog)
- We use this on our team to enable performant aggregation of complicated statistics without storing all rows in Spark memory.

```
object DoubleSketch {
  val k: Int = 512
  def apply(serialized: Array[Byte]): DoublesSketch = DoublesSketch.wrap(Memory.wrap(serialized))
  def sketch(data: Array[Double]): DoublesSketch = sketch(k, data)
  def sketch(k: Int, data: Array[Double]): DoublesSketch = {
    val s = DoublesSketch.builder().setK(k).build()
   data.foreach { s.update }
  }
  def union(k: Int, a: Array[Double], b: Array[Double]): DoublesSketch = {
    val u = DoublesUnion.builder().setMaxK(k).build()
    u.update(sketch(k, a))
    u.update(sketch(k, b))
    u.getResult
  def union(k: Int, a: DoublesSketch, b: DoublesSketch): DoublesSketch = {
    val u = DoublesUnion.builder().setMaxK(k).build()
    u.update(a)
    u.update(b)
    u.getResult
```



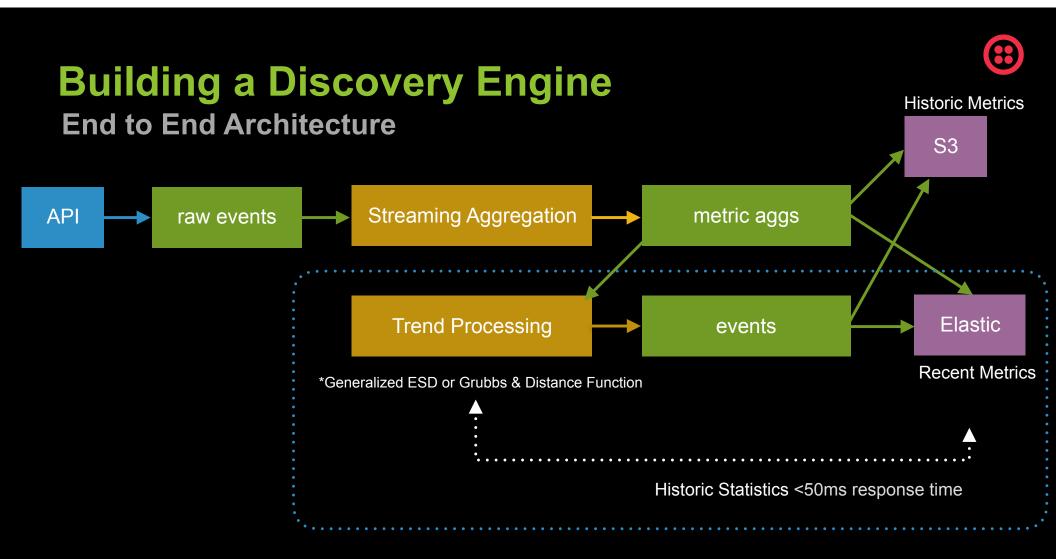


- Testing should feel good.
- Ensure you solve Encoder issues early
- Make sure you don't add regression's into your 24/7/365 apps!

```
test("Should aggregate call events") {
 import session.implicits._
 val appConfig = appConfigForTest()
 val scenario = TestHelper.loadScenario[CallEvent](s"$pathToTestScenarios/pdd_events.json")
 val scenarioIter = scenario.toIterator
 scenario.nonEmptv shouldBe true
 testUtils.createTopic(kafkaTopic, partitions, overwrite = true)
 sendNextMessages(scenarioIter, 30, _.getEventId, _.getLoggedEventTime)
 val trendDiscoveryApp = new TrendDiscoveryApp(appConfigForTest(), session)
 val eventAggregation = EventAggregation(appConfig)
 eventAggregation.process(trendDiscoveryApp.readKafkaStream())(session)
   .writeStream
    .quervName("calleventaggs")
   .format("memory")
   .outputMode(eventAggregation.outputMode)
   .start()
   .processAllAvailable()
 val res = session
   .sql("select avg(stats.p99) from calleventaggs")
   .collect()
    .map { r =>
     r.getAs[Double](0) }
   .head
 DiscoveryUtils.round(res) shouldEqual 7.13
```



#DDSAIS16







**Anatomy of the application: Streaming Aggregation Query** 

Building the Continuous Trend Detector

- 1. Read from Kafka (raw events)
- 2. Look up \*historic statistics from Elastic Search or Redis
- 3. Generate an anomaly stream based off of distance to statistical norm
- 4. Emit to Kafka for downstream use





Discovering trends in massive data sets doesn't come for free. There are core principals and ideologies that must be followed in order to extract interesting patterns from your data.





#### Questions? / Thanks



shaines@twilio.com

Thanks to Marek Radonsky and Rohit Gupta from the Voice Insights team who helped make these patterns possible.

